

Analytics 511: Homework #11

Hanjing Wang

```
set.seed(1)
```

```
Hitters = read.csv('Hitters.csv')
```

1

(1)

```
library("glmnet")
```

```
## Warning: package 'glmnet' was built under R version 3.4.4
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Warning: package 'foreach' was built under R version 3.4.4
```

```
## Loaded glmnet 2.0-13
```

```
Hitters = na.omit(Hitters)
```

We use LASSO regression to predict Salary from the other numeric predictors.


```
matrix <- model.matrix(lm(Salary ~ .-X-League-Division-NewLeague, data = Hitters))
```

```
fit.lasso = glmnet(matrix, Hitters$Salary, alpha = 1)
```

```
# plot the coefficient trajectories as a function of the regularization parameter lambda
```

```
plot(fit.lasso, xvar = "lambda", lwd = 1, label=TRUE)
```

```
grid(col = 1)
```



```
# we could also show the coefficients when there are only three variables left
fit.lasso$beta[,5]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
##  0.00000000  0.00000000  0.08064999  0.00000000  0.00000000  0.00000000
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##  0.00000000  0.00000000  0.00000000  0.00000000  0.00000000  0.06719193
##      CRBI      CWalks      PutOuts      Assists      Errors
##  0.17823025  0.00000000  0.00000000  0.00000000  0.00000000
```

From the plot, we could see that the final three predictors that remain in the model are Hits, CRuns, CRBI.

Then, We use cross-validation to find the optimal value of the regularization penalty.

```
cv.lasso <- cv.glmnet(matrix, Hitters$Salary, alpha = 1)
bestlambda<- cv.lasso$lambda.min
bestlambda
```

```
## [1] 2.935124
```

Finally, we use the bestlambda to fit the model.

```
fit.lasso1 = glmnet(matrix, Hitters$Salary, alpha = 1, lambda = bestlambda)
fit.lasso1$beta
```

```
## 17 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept)  .
## AtBat      -1.7099436
## Hits       6.0823207
```

```
## HmRun      .
## Runs       .
## RBI        0.2602043
## Walks      4.7542167
## Years      -9.0357077
## CAtBat     .
## CHits      .
## CHmRun     0.5745583
## CRuns      0.6973902
## CRBI       0.3033389
## CWalks     -0.4924500
## PutOuts    0.2804673
## Assists    0.1839622
## Errors     -1.7085130
```


We could see that 12 predictors are left in that model.

(2)

Repeat with Ridge Regression.

```
matrix1 <- model.matrix(lm(Salary ~ .-X-League-Division-NewLeague, data =Hitters))
fit.ridge = glmnet(matrix1, Hitters$Salary, alpha = 0)

# plot the coefficient trajectories as a function of the regularization parameter lambda
plot(fit.ridge, xvar = "lambda", lwd = 1,label=TRUE)
grid(col = 1)
```



```
cv.ridge <- cv.glmnet(matrix1, Hitters$Salary, alpha = 0)
bestlambda<- cv.ridge$lambda.min
bestlambda
```

```
## [1] 28.01718
```

```
# the coefficients
```

```
fit.ridge1 = glmnet(matrix, Hitters$Salary, alpha = 0, lambda = bestlambda)
fit.ridge1$beta
```

```
## 17 x 1 sparse Matrix of class "dgCMatrix"
```

```
##                s0
```

```
## (Intercept)  .
```

```
## AtBat      -0.723504487
```

```
## Hits       2.704625078
```

```
## HmRun      -2.595387303
```

```
## Runs       1.351260127
```

```
## RBI        1.226757322
```

```
## Walks      3.254158746
```

```
## Years     -8.485767019
```

```
## CAtBat     -0.001692903
```

```
## CHits      0.125773172
```

```
## CHmRun     0.661490023
```

```
## CRuns      0.297444435
```

```
## CRBI       0.242310378
```

```
## CWalks    -0.228174794
```

```
## PutOuts    0.270173394
```

```
## Assists      0.169600493
## Errors      -3.413387246
```

2

The bias-variance trade-off is that in a series of models, the models which have a lower bias for the parameter estimation will have a higher variance, and vice versa. The regularization will shrink the coefficient estimates towards zero. Shrinking the coefficient estimate will lead to a substantial reduction in the variance of the predictions, at the expense of a slight increase in bias so that the overall prediction accuracy will be improved. For example, in 1.1, the lasso regression model using the best lambda generating from the cross validation will have a slightly higher bias than the model using all the predictors but also will have a much lower variance. Compared the model where there are only three predictors left, the model with the best lambda will have a lower bias but higher variance.